

Combining TIFF, HDF5, and Zarr into a Single Image File Format

Mark Kittisopikul

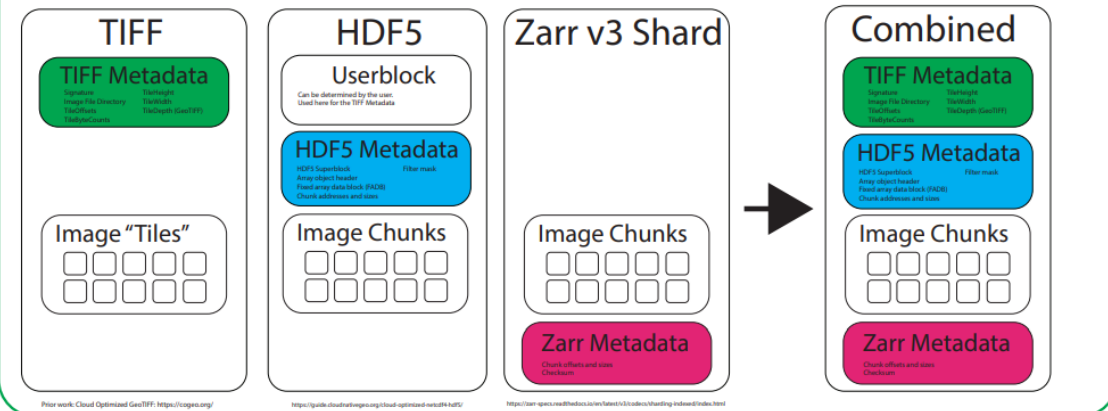
October 3, 2025

1 Abstract

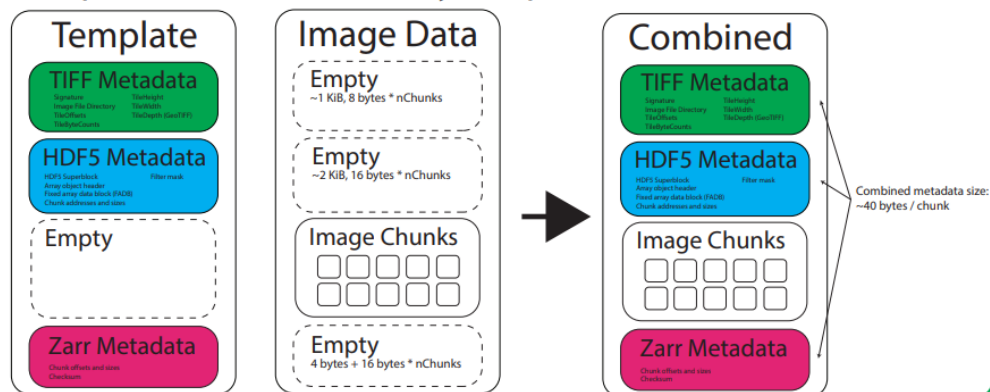


“Situation: There are 3 competing standards” the last card of a popular XKCD cartoon (#927) might read if applied to recent microscopy bioimaging formats. TIFF, HDF5, and Zarr have all been used to store images as part of popular standards and formats (OME-TIFF, BDV-HDF5, OME-Zarr). The cartoon humorously points out the tendency to create new standards while discounting prior efforts. To combat the proliferation of formats I examine similarities between TIFF, HDF5, and Zarr shard containers. I then exploit them to create a combined data container that is simultaneously a TIFF file, a HDF5 file, and a Zarr version 3 shard without duplicating the image pixel or volumetric voxel data. This combined format is compatible with multiple viewers and image analysis pipelines. Additionally, the techniques involved provide a path to convert between the formats with minimal processing or overhead. In practice, the combined format avoids redundant copies of data while providing great utility to the microscope user. The combined format is a great candidate for a microscope acquisition format as it satisfies both short term needs to view microscope output in traditional viewers while integrating into next generation image analysis pipelines.

Cloud-optimized file formats have aggregated metadata, allowing for a combined file format



A simple implementation via binary templates



Implementation Outline

LibTIFF

1. Write TIFF signature and first image file directory (IFD)
2. Close the TIFF file
3. Manually pad the file to the combined size of the TIFF and HDF5 headers (~3 KiB)
4. Reopen the TIFF file with LibTIFF and write some placeholder pixels.

LibHDF5

1. Create a HDF5 file with the following settings:
 - A. Library version set to v1.12
 - B. A user block of ~1 KiB for TIFF header
 - C. A meta block size large enough to fit all HDF5 metadata (2 KiB)
2. Write a chunked dataset (an array) with the following settings:
 - A. Allocation time set to "eager"
 - B. Fill time set to "never"
3. Write a contiguous dataset for the zarr index.

Zarr

1. Read the chunk / tile locations from HDF5 or TIFF
2. Write a Zarr shard chunk index and checksum as a HDF5 dataset

chunk (0, 0)	chunk (0, 1)	chunk (1, 0)	chunk (1, 1)	checksum
offset nbytes	offset nbytes	offset nbytes	offset nbytes	checksum
uint64 uint64	uint64 uint64	uint64 uint64	uint64 uint64	uint32

1.1 1. Write the Combined File Format

```
[1]: import header_formats
```

```
[2]: import os
import shutil
import numpy as np
def run_demo():
    # create a template file
```

```

header_formats.tiff_hdf5_zarr("test.tiff.hdf5.zarr", (256,256), "uint16",
↪chunks=(128,128))
header_formats.create_zarr3("demo/test.zarr")

# extract header and footer from template
header, footer = header_formats.read_header_footer("test.tiff.hdf5.zarr")

# create chunks
A = np.full((128, 128), 0, dtype="uint16")
B = np.full((128, 128), 2**14-2, dtype="uint16")
C = np.full((128, 128), 2*2**14-2, dtype="uint16")
D = np.full((128, 128), 3*2**14-2, dtype="uint16")

os.makedirs("demo/test.zarr/c/0/")

# write header, tiles and footer to demo file
with open("demo/demo.hdf5.zarr.tiff", "wb") as f:
    f.write(header)
    f.write(A)
    f.write(B)
    f.write(C)
    f.write(D)
    f.write(footer)

# copy one file to many files with different file extensions
# TIFF
shutil.copyfile("demo/demo.hdf5.zarr.tiff", "demo/demo.tiff")
# HDF5
shutil.copyfile("demo/demo.hdf5.zarr.tiff", "demo/demo.h5")
# Zarr v3 shard
shutil.copyfile("demo/demo.hdf5.zarr.tiff", "demo/test.zarr/c/0/0")

```

```
[3]: run_demo()
```

MissingRequired: TIFF directory is missing required "TileOffsets" field.
MissingRequired: TIFF directory is missing required "TileOffsets" field.

1.2 2. Check the HDF5 and TIFF File Structure

```
[4]: !pixi run h5ls -va demo/demo.hdf5.zarr.tiff
```

```

Opened "demo/demo.hdf5.zarr.tiff" with sec2 driver.
data                                Dataset {256/256, 256/256}
  Location:    1:195
  Links:       1
  Modified:    2025-10-03 17:39:17 EDT
  Chunks:      {128, 128} 32768 bytes
  Storage:     131072 logical bytes, 131072 allocated bytes, 100.00% utilization

```

```

Type:      native unsigned short
Address: 2048
      Flags   Bytes   Address   Logical Offset
      =====
0x00000000   32768     2048 [0, 0, 0]
0x00000000   32768    34816 [0, 128, 0]
0x00000000   32768    67584 [128, 0, 0]
0x00000000   32768   100352 [128, 128, 0]
zarrindex      Dataset {68/68}
  Location:   1:479
  Links:      1
  Storage:    68 logical bytes, 68 allocated bytes, 100.00% utilization
  Type:       native unsigned char
  Address:    133120

```

```
[5]: !pixi run tiffinfo -s demo/demo.hdf5.zarr.tiff
```

```

=== TIFF directory 0 ===vironment
TIFF Directory at offset 0x86 (134)
  Image Width: 256 Image Length: 256
  Tile Width: 128 Tile Length: 128
  Bits/Sample: 16
  Sample Format: unsigned integer
  Compression Scheme: None
  Photometric Interpretation: min-is-black
  Orientation: row 0 top, col 0 lhs
  Planar Configuration: single image plane
  4 Tiles:
    0: [   3072,   32768]
    1: [  35840,   32768]
    2: [   68608,   32768]
    3: [ 101376,   32768]

```

1.3 3. Read the Data using h5py, libtiff, and tensorstore

```
[6]: import h5py
with h5py.File("demo/demo.hdf5.zarr.tiff") as h5f:
    h5data = h5f["data"][:]

h5data
```

```
[6]: array([[ 0,  0,  0, ..., 16382, 16382, 16382],
 [ 0,  0,  0, ..., 16382, 16382, 16382],
 [ 0,  0,  0, ..., 16382, 16382, 16382],
 ...,
 [32766, 32766, 32766, ..., 49150, 49150, 49150],
 [32766, 32766, 32766, ..., 49150, 49150, 49150],
```

```
[32766, 32766, 32766, ..., 49150, 49150, 49150]],  
shape=(256, 256), dtype=uint16)
```

```
[7]: from libtiff import TIFF  
tif = TIFF.open("demo/demo.hdf5.zarr.tiff", "r")  
tiff_data = tif.read_image()  
tif.close()  
tiff_data
```

```
[7]: array([[ 0, 0, 0, ..., 16382, 16382, 16382],  
[ 0, 0, 0, ..., 16382, 16382, 16382],  
[ 0, 0, 0, ..., 16382, 16382, 16382],  
...,  
[32766, 32766, 32766, ..., 49150, 49150, 49150],  
[32766, 32766, 32766, ..., 49150, 49150, 49150],  
[32766, 32766, 32766, ..., 49150, 49150, 49150]],  
shape=(256, 256), dtype=uint16)
```

```
[8]: import tensorstore as ts  
ts.open({  
    "driver": "zarr3",  
    "kvstore": {  
        "driver": "file",  
        "path": "demo/test.zarr/"  
    },  
}).result().read().result()
```

```
[8]: array([[ 0, 0, 0, ..., 16382, 16382, 16382],  
[ 0, 0, 0, ..., 16382, 16382, 16382],  
[ 0, 0, 0, ..., 16382, 16382, 16382],  
...,  
[32766, 32766, 32766, ..., 49150, 49150, 49150],  
[32766, 32766, 32766, ..., 49150, 49150, 49150],  
[32766, 32766, 32766, ..., 49150, 49150, 49150]],  
shape=(256, 256), dtype=uint16)
```

```
[9]: import h5py  
with h5py.File("demo/demo.hdf5.zarr.tiff") as h5f:  
    h5data = h5f["data"][:]  
    h5zarrindex = h5f["zarrindex"][:]  
    offsets_and_bytes = header_formats.  
        ↪get_hdf5_chunk_offsets_and_bytes(h5f["data"])  
  
h5zarrindex[0:-4].view(np.uint64)  
offsets_and_bytes
```

```
[9]: array([[ 3072, 32768],
           [ 35840, 32768],
           [ 68608, 32768],
           [101376, 32768]], dtype=uint64)
```

1.4 4. Rewrite the data using h5py

```
[10]: import h5py
import shutil
with h5py.File("demo/demo.hdf5.zarr.tiff", "r+") as h5f:
    h5f["data"][:128,:128] = 1
    h5f["data"][:128,128:] = 2
    h5f["data"][128:,:128] = 3
    h5f["data"][128:,128:] = 4

# copy to the zarr shard, consider a symlink on Linux systems
shutil.copyfile("demo/demo.hdf5.zarr.tiff", "demo/test.zarr/c/0/0")
```

```
[10]: 'demo/test.zarr/c/0/0'
```

1.5 5. Read the updated data using h5py, libtiff, and tensorstore

```
[11]: with h5py.File("demo/demo.hdf5.zarr.tiff") as h5f:
    print(h5f["data"][:])
```

```
[[1 1 1 ... 2 2 2]
 [1 1 1 ... 2 2 2]
 [1 1 1 ... 2 2 2]
 ...
 [3 3 3 ... 4 4 4]
 [3 3 3 ... 4 4 4]
 [3 3 3 ... 4 4 4]]
```

```
[12]: tif = TIFF.open("demo/demo.hdf5.zarr.tiff", "r")
print(tif.read_image())
tif.close()
```

```
[[1 1 1 ... 2 2 2]
 [1 1 1 ... 2 2 2]
 [1 1 1 ... 2 2 2]
 ...
 [3 3 3 ... 4 4 4]
 [3 3 3 ... 4 4 4]
 [3 3 3 ... 4 4 4]]
```

```
[13]: ts.open({
    "driver": "zarr3",
    "kvstore": {
```

```
        "driver": "file",
        "path": "demo/test.zarr/"
    },
}).result().read().result()
```

```
[13]: array([[1, 1, 1, ..., 2, 2, 2],
            [1, 1, 1, ..., 2, 2, 2],
            [1, 1, 1, ..., 2, 2, 2],
            ...,
            [3, 3, 3, ..., 4, 4, 4],
            [3, 3, 3, ..., 4, 4, 4],
            [3, 3, 3, ..., 4, 4, 4]], shape=(256, 256), dtype=uint16)
```